

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)

data = np.hstack([raw_df.values[:, :-2], raw_df.values[:, -2]])
target = raw_df.values[:, -2]

columns = [
    "CRIM", "ZN", "INDUS", "CHAS", "NOX", "RM", "AGE",
    "DIS", "RAD", "TAX", "PTRATIO", "B", "LSTAT"
]

df = pd.DataFrame(data, columns=columns)
df['MEDV'] = target


df.head()

{
  "summary": {
    "name": "df",
    "rows": 506,
    "fields": [
      {
        "column": "CRIM",
        "properties": {
          "dtype": "number",
          "std": 8.601545105332487,
          "min": 0.00632,
          "max": 88.9762,
          "num_unique_values": 504,
          "samples": [
            0.09178, 0.05644, 0.10574
          ],
          "semantic_type": "\",
          "description": "\n"
        }
      },
      {
        "column": "ZN",
        "properties": {
          "dtype": "number",
          "std": 23.322452994515036,
          "min": 0.0,
          "max": 100.0,
          "num_unique_values": 26,
          "samples": [
            25.0, 30.0, 18.0
          ],
          "semantic_type": "\",
          "description": "\n"
        }
      },
      {
        "column": "INDUS",
        "properties": {
          "dtype": "number",
          "std": 6.8603529408975845,
          "min": 0.46,
          "max": 27.74,
          "num_unique_values": 76,
          "samples": [
            8.14, 1.47, 1.22
          ],
          "semantic_type": "\",
          "description": "\n"
        }
      },
      {
        "column": "CHAS",
        "properties": {
          "dtype": "number",
          "std": 0.2539940413404118,
          "min": 0.0,
          "max": 1.0,
          "num_unique_values": 2,
          "samples": [
            1.0, 0.0
          ],
          "semantic_type": "\",
          "description": "\n"
        }
      }
    ]
  }
}

```

```
    },\n      {\n        \"column\": \"NOX\", \n        \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 0.11587767566755611, \n          \"min\": 0.385, \n          \"max\": 0.871, \n          \"num_unique_values\": 81, \n          \"samples\": [\n            0.401, \n            0.538\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n        }, \n        {\n          \"column\": \"RM\", \n          \"properties\": {\n            \"dtype\": \"number\", \n            \"std\": 0.7026171434153237, \n            \"min\": 3.561, \n            \"max\": 8.78, \n            \"num_unique_values\": 446, \n            \"samples\": [\n              6.849, \n              4.88\n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\"\n          }, \n          {\n            \"column\": \"AGE\", \n            \"properties\": {\n              \"dtype\": \"number\", \n              \"std\": 28.148861406903638, \n              \"min\": 2.9, \n              \"max\": 100.0, \n              \"num_unique_values\": 356, \n              \"samples\": [\n                51.8, \n                33.8\n              ], \n              \"semantic_type\": \"\", \n              \"description\": \"\"\n            }, \n            {\n              \"column\": \"DIS\", \n              \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 2.1057101266276104, \n                \"min\": 1.1296, \n                \"max\": 12.1265, \n                \"num_unique_values\": 412, \n                \"samples\": [\n                  2.2955, \n                  4.2515\n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\"\n              }, \n              {\n                \"column\": \"RAD\", \n                \"properties\": {\n                  \"dtype\": \"number\", \n                  \"std\": 8.707259384239377, \n                  \"min\": 1.0, \n                  \"max\": 24.0, \n                  \"num_unique_values\": 9, \n                  \"samples\": [\n                    7.0, \n                    2.0\n                  ], \n                  \"semantic_type\": \"\", \n                  \"description\": \"\"\n                }, \n                {\n                  \"column\": \"TAX\", \n                  \"properties\": {\n                    \"dtype\": \"number\", \n                    \"std\": 168.53711605495926, \n                    \"min\": 187.0, \n                    \"max\": 711.0, \n                    \"num_unique_values\": 66, \n                    \"samples\": [\n                      370.0, \n                      666.0\n                    ], \n                    \"semantic_type\": \"\", \n                    \"description\": \"\"\n                  }, \n                  {\n                    \"column\": \"PTRATIO\", \n                    \"properties\": {\n                      \"dtype\": \"number\", \n                      \"std\": 2.164945523714446, \n                      \"min\": 12.6, \n                      \"max\": 22.0, \n                      \"num_unique_values\": 46, \n                      \"samples\": [\n                        19.6, \n                        15.6\n                      ], \n                      \"semantic_type\": \"\", \n                      \"description\": \"\"\n                    }, \n                    {\n                      \"column\": \"B\", \n                      \"properties\": {\n                        \"dtype\": \"number\", \n                        \"std\": 91.29486438415779, \n                        \"min\": 0.32, \n                        \"max\": 396.9, \n                        \"num_unique_values\": 357, \n                        \"samples\": [\n                          396.24, \n                          395.11\n                        ], \n                        \"semantic_type\": \"\", \n                        \"description\": \"\"\n                      }, \n                      {\n                        \"column\": \"LSTAT\", \n                        \"properties\": {\n                          \"dtype\": \"number\", \n                          \"std\": 7.141061511348571, \n                          \"min\": 1.73, \n                          \"max\": 37.97, \n                          \"num_unique_values\": 455, \n                          \"samples\": [\n                            6.15, \n                            4.32\n                          ], \n                          \"semantic_type\": \"\", \n                          \"description\": \"\"\n                        }\n                      }\n                    }\n                  }\n                }\n              }\n            }\n          }\n        }\n      }\n    }\n  }\n}\n
```

```
n      },\n      {\n        \"column\": \"MEDV\", \n        \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 9.19710408737982, \n          \"min\": 5.0, \n          \"max\": 50.0, \n          \"num_unique_values\": 229, \n          \"samples\": [\n            14.1, \n            22.5\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n        }\n      }\n    ]\n  },\n  \"type\": \"dataframe\", \n  \"variable_name\": \"df\"\n}\n\nX = df.drop('MEDV', axis=1)\ny = df['MEDV']\n\nX_train, X_test, y_train, y_test = train_test_split(X, y,\ntest_size=0.2, random_state=42)\n\nmodel = LinearRegression()\nmodel.fit(X_train, y_train)\n\nLinearRegression()\n\ny_pred = model.predict(X_test)\n\nprint("Mean Squared Error:", mean_squared_error(y_test, y_pred))\nprint("R2 Score:", r2_score(y_test, y_pred))\n\nMean Squared Error: 24.291119474973478\nR2 Score: 0.6687594935356326\n\nsns.scatterplot(x=df['RM'], y=df['MEDV'])\nplt.title("Average Rooms vs Median House Price")\nplt.xlabel("RM (Avg Rooms)")\nplt.ylabel("MEDV ($1000s)")\nplt.grid(True)\nplt.show()
```



```

from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

iris = load_iris()
X = iris.data
y = iris.target

df = pd.DataFrame(X, columns=iris.feature_names)
df['species'] = y

print(df.isnull().sum())
    sepal length (cm)      0
    sepal width (cm)       0
    petal length (cm)      0

```

```
petal width (cm)      0
species              0
dtype: int64

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

model = LogisticRegression(max_iter=200)
model.fit(X_train, y_train)


LogisticRegression(max_iter=200)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

Accuracy: 1.0

cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, cmap="Blues", fmt='d',
            xticklabels=iris.target_names,
            yticklabels=iris.target_names)
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



```

feature_importance = pd.Series(
    abs(model.coef_).mean(axis=0),
    index=iris.feature_names
).sort_values(ascending=False)

print("Feature Importance:\n", feature_importance)

Feature Importance:
petal length (cm)      1.725414
petal width (cm)       1.183001
sepal width (cm)       0.641724
sepal length (cm)      0.338936
dtype: float64

from sklearn.datasets import load_iris
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns

```

```


iris = load_iris()
X = pd.DataFrame(iris.data, columns=iris.feature_names)
y = iris.target

inertia = []
k_range = range(1, 11)

for k in k_range:
    model = KMeans(n_clusters=k, random_state=42)
    model.fit(X)
    inertia.append(model.inertia_)

plt.plot(k_range, inertia, marker='o')
plt.title('Elbow Method for Optimal k')
plt.xlabel('Number of clusters (k)')
plt.ylabel('Inertia (WCSS)')
plt.grid(True)
plt.show()

```



```

kmeans = KMeans(n_clusters=3, random_state=42)
X['cluster'] = kmeans.fit_predict(X)


```

```

pca = PCA(n_components=2)
X_pca = pca.fit_transform(iris.data)

plt.figure(figsize=(8, 6))
sns.scatterplot(x=X_pca[:, 0], y=X_pca[:, 1], hue=X['cluster'],
palette='Set1')
plt.title("K-Means Clusters (PCA 2D Projection)")
plt.xlabel("PCA Component 1")
plt.ylabel("PCA Component 2")
plt.legend(title='Cluster')
plt.show()

```



```

from sklearn.metrics import adjusted_rand_score

ari = adjusted_rand_score(y, X['cluster'])
print("Adjusted Rand Index (ARI):", ari)

Adjusted Rand Index (ARI): 0.7163421126838476

```

*#A: 3 clusters*

*#B*